

Literature Review: Chess Programming Techniques

Comprehensive Citation Analysis

Citation	Key Contributions	Relationship to Other Works	Role in Review
Shannon (1950) <i>Programming a Computer for Playing Chess</i>	<ul style="list-style-type: none"> Frames chess as computational problem Identifies 10^{120} position space impossibility Proposes evaluation function $f(P)$ Introduces Type A/B search strategies Recognizes horizon effect Suggests material, mobility, pawn structure, king safety as evaluation factors 	<ul style="list-style-type: none"> Foundation for all subsequent work Directly influences Knuth & Moore's formalization Evaluation principles adopted by all HCE engines Type B strategy leads to selective extensions Basically is the foundational paper that drove all of chess programming 	Establishes theoretical foundations; defines core problems (intractable search space, need for evaluation, horizon effect) that motivate decades of innovation
Knuth & Moore (1975) <i>Analysis of Alpha-Beta Pruning</i>	<ul style="list-style-type: none"> Formalizes minimax as $F(p), G(p)$ Proves alpha-beta optimality Demonstrates deep cutoffs advantage Shows F_3 improvements impossible Best case: $W^{\frac{D}{2}} + W^{\frac{D}{2}} - 1$ nodes 	<ul style="list-style-type: none"> Builds on Shannon's minimax concept Validates alpha-beta vs. all alternatives Referenced by Marsland, Brange, all modern pruning work Later challenged by AlphaZero's MCTS success 	Provides mathematical proof of alpha-beta's theoretical optimality within minimax framework; establishes performance bounds guiding implementation decisions
Björnsson & Marsland (2000) <i>Review of Game-Tree Pruning</i>	<ul style="list-style-type: none"> Comprehensive pruning taxonomy TPT essential for pruning & move ordering Null-move, futility pruning techniques Quiescence search importance Move ordering heuristics (killer, history) 	<ul style="list-style-type: none"> Synthesizes Knuth & Moore theory with practice Connects to Shannon's Type B strategy Foundation for Brange, Tesseract implementations Complements Zobrist's hashing 	Bridges theory and practice; comprehensive survey of pruning landscape from foundational to advanced techniques
Zobrist (1970) <i>New Hashing Method</i>	<ul style="list-style-type: none"> Incremental hashing via XOR 2-4 XOR operations per update 0.000003% collision rate with 1B positions Enables efficient transposition tables 	<ul style="list-style-type: none"> Enables Marsland's TPT recommendations Universal adoption (Stockfish, LC0, all engines) Complements magic bitboards for efficiency Used with all search techniques Now is the standard for hashing 	Solves position hashing problem; enables memory-aided search techniques critical to modern engine performance

Citation	Key Contributions	Relationship to Other Works	Role in Review
Nasu (2018) <i>NNUE: Efficiently Updatable Neural Networks</i>	<ul style="list-style-type: none"> Introduces NNUE architecture Incremental accumulator mechanism HalfKP feature representation CPU-optimized neural evaluation Quantized integer arithmetic (8-16 bit) 	<ul style="list-style-type: none"> Bridges HCE limitations with neural learning Alternative to AlphaZero's GPU approach Integrated into Stockfish (2020) Complements alpha-beta search 	Enables neural evaluation in traditional engines; represents hybrid paradigm combining alpha-beta efficiency with learned patterns
Silver et al. (2017) <i>Mastering Chess and Shogi (AlphaZero)</i>	<ul style="list-style-type: none"> MCTS + deep neural networks Self-play reinforcement learning 80K pos/s vs Stockfish's 70M pos/s Defeats Stockfish 28-0-72 despite 1000x fewer nodes Policy + value network architecture 	<ul style="list-style-type: none"> Challenges Knuth & Moore's alpha-beta dominance Contrasts with Shannon's HCE approach Influences Leela Chess Zero development Motivates NNUE as CPU alternative 	Demonstrates paradigm shift from exhaustive to selective search; proves learned evaluation can surpass handcrafted heuristics
Fiekas (2018) <i>Finding Hash Functions for Bitboard Move Generation</i>	<ul style="list-style-type: none"> Magic bitboard hash functions Multiplicative hashing technique Brute-force magic number discovery Constant-time move generation PEXT comparison shows only 2.3% speedup 	<ul style="list-style-type: none"> Extends bitboard work (1970s origins) Compared with BMI2 PEXT instructions Used by Stockfish, all modern engines Dives deeper into magic bitboards 	Solves sliding piece move generation; provides constant-time lookup essential for competitive performance
de Groot (1965) <i>Thought and Choice in Chess</i>	<ul style="list-style-type: none"> Coins "progressive deepening" term Studies chess cognition Iterative search concept Human pattern recognition insights 	<ul style="list-style-type: none"> Theoretical foundation for iterative deepening Influences Marsland's search enhancements Applied by all modern engines Cognitive basis for search strategies 	Introduces iterative deepening concept; provides cognitive basis for human-inspired search strategies
Bijl & Tiet (2021) <i>Exploring Modern Chess Engine Architectures</i>	<ul style="list-style-type: none"> Empirical architecture comparison 0x88 nearly matches bitboards in Perft (46.5 vs 48.8 Mn/s) Bitboards dominate in evaluation, not move generation Move generation only 10% of engine time 	<ul style="list-style-type: none"> Challenges conventional bitboard justification Validates magic bitboards vs PEXT findings Tests techniques from Marsland, Fiekas Supports NNUE need for learned evaluation 	Provides empirical validation of architectural choices; reveals nuanced performance tradeoffs challenging conventional assumptions

Citation	Key Contributions	Relationship to Other Works	Role in Review
	<ul style="list-style-type: none"> Parameter tuning improves 15% win rate 		
Brange (2021) <i>Evaluating Heuristic and Algorithmic Improvements (KLAS)</i>	<ul style="list-style-type: none"> KLAS engine case study MVV-LVA: 68.5% execution time reduction Iterative deepening + PV: 28.7% faster Lazy SMP: 33.1% speedup (4 cores) Empirical technique validation 	<ul style="list-style-type: none"> Implements Marsland's techniques Validates Zobrist, magic bitboards Compares YBWC vs Lazy SMP Supports Tesseract findings 	Provides quantitative performance impacts; demonstrates real-world effectiveness of theoretical techniques
Vrzina (2023) <i>Piece By Piece Building a Strong Engine (Tesseract)</i>	<ul style="list-style-type: none"> Tesseract engine development journey Quiescence Search: 62.5% faster, +1206 eval score PSTs: biggest eval impact (5640→8255) 16-bit move encoding: 50% faster generation Lazy SMP: 40% speedup LMR challenges (score drop 8584→8124) 	<ul style="list-style-type: none"> Validates techniques from Marsland, Brange Confirms bitboard + mailbox hybrid benefits LMR challenges align with AlphaDeepChess Supports Shannon's evaluation factors 	Documents practical implementation journey; reveals implementation challenges (LMR sensitivity) and empirical performance gains
Columbia et al. (2023) <i>Chess Move Generation Using Bitboards</i>	<ul style="list-style-type: none"> Bitboard move generation analysis Legal vs pseudo-legal comparison Lookup table techniques Modern engine preferences documented 	<ul style="list-style-type: none"> Extends 1970s bitboard origins Complements Fiekas's magic bitboards Aligns with Stockfish implementation choices 	Clarifies move generation approaches; explains pseudo-legal preference in modern competitive engines
Rasmussen (2004) <i>Parallel Chess Searching and Bitboards</i>	<ul style="list-style-type: none"> Parallel search fundamental challenges Young Brothers Wait Concept (YBWC) algorithm Lazy SMP concept 9.2× speedup on 22 processors (not 22×) Synchronization overhead analysis Tree size growth roughly linear 	<ul style="list-style-type: none"> Theoretical foundation for Brange, Vrzina work Explains Stockfish's YBWC→Lazy SMP switch Connects to alpha-beta's sequential nature Identifies parallelization limits 	Identifies fundamental parallelization challenges; explains sublinear scaling and inherent tradeoffs in parallel search
Girón & Wang (2025) <i>AlphaDeepChess: Alpha-Beta Engine</i>	<ul style="list-style-type: none"> AlphaDeepChess engine development LMR implementation challenges and failures YBWC performance degradation 	<ul style="list-style-type: none"> Validates Rasmussen's YBWC concerns Confirms LMR sensitivity (also in Tesseract) Contrasts with Stockfish's successful implementation 	Highlights implementation difficulty gap; shows advanced techniques require sophisticated supporting systems to succeed

Citation	Key Contributions	Relationship to Other Works	Role in Review
	<ul style="list-style-type: none"> Synchronization overhead issues Move ordering criticality demonstrated 	<ul style="list-style-type: none"> Highlights implementation difficulty 	
Świechowski et al. (2022) <i>Monte Carlo Tree Search: A Review</i>	<ul style="list-style-type: none"> MCTS comprehensive review UCB1 selection strategy MCTS variants and modifications Application across domains Pre-neural MCTS performance 	<ul style="list-style-type: none"> Contextualizes AlphaZero's MCTS approach Documents historical MCTS weakness vs alpha-beta Explains neural networks as game-changer Provides theoretical background 	Provides MCTS theoretical background; explains why neural guidance was necessary for MCTS superiority over alpha-beta
Stockfish Team (2025) <i>NNUE Documentation</i>	<ul style="list-style-type: none"> NNUE implementation in Stockfish HalfKAv2 architecture details Accumulator mechanism explained Integration with alpha-beta search 100+ Elo gain documented 	<ul style="list-style-type: none"> Implements Nasu's NNUE architecture Alternative to AlphaZero's GPU approach Combines with traditional alpha-beta techniques Represents current state-of-the-art 	Documents successful hybrid approach at highest competitive level; proves neural evaluation compatible with classical search
Plaat (1997) <i>A Minimax Algorithm Faster than NegaScout</i>	<ul style="list-style-type: none"> MTD(f) algorithm (Memory-enhanced Test Driver) Multiple minimal window searches Convergence to minimax value Strong transposition table dependency Alternative to PVS/ NegaScout 	<ul style="list-style-type: none"> Alternative to PVS/ NegaScout approach Builds on Knuth & Moore's alpha-beta bounds Less widely adopted than PVS due to TPT dependency Requires strong hashing infrastructure 	Illustrates continued algorithmic refinement; shows innovation continues despite theoretical optimality proofs of alpha-beta

Table 1: Analysis of Sources based on their contributions